

Yang Yiming

Gender: Male Objective: AI Agent Engineer Location: Hangzhou

Phone: 17369709317 Email: yangyiming6372@gmail.com GitHub: <https://github.com/yangyiming586>

Education

2011.09 - 2015.06 South China University of Technology Artificial Intelligence (B.S.)

Courses: Software Engineering, Computer Organization, Compilers, Operating Systems, Distributed Systems, Design Patterns

Summary

Backend engineer building production AI-agent workflows, RAG pipelines and runtime observability, skilled in multi-agent orchestration and tool governance.

Work Experience

2018.07 - Present Xiaohongshu Senior AI Agent Engineer

- Added structured constraints and validation on LLM output to reduce hallucination, raising key-field extraction accuracy above 90%.
- Integrated an MCP resume-evidence service and dynamic skill loading for evidence synthesis and tool governance, ensuring online service stability.
- Wired the full LLM call chain into Langfuse, tracking hit count, top score, fallback rate and latency per step, significantly lowering maintenance cost.
- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 38%.
- Built an offline evaluation set and automated scoring covering faithfulness, relevance and usability, significantly lowering maintenance cost.

2015.07 - 2018.06 Li Auto AI Agent Engineer

- Implemented hybrid RAG fusing BM25 lexical recall, vector search, reranking and fallback visibility, raising evidence hit rate by 42%.
- Built backend of a multi-agent resume evaluation workflow chaining intent router, parser, JD matcher, technical/risk evaluators and report generator.
- Handled 4960K agent invocations per day while keeping the online service stable, and related standards were adopted by the team.
- Tuned Milvus index and sharding strategy, reducing retrieval latency from 1800ms to 180ms, with experience documented and shared internally.

Projects

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Supported multi-turn conversation memory and context management.
- Added reranking and citation tracing to reduce off-topic answers.
- Implemented document chunking and multi-route recall to improve long-document QA.
- Built online evaluation and a feedback loop.

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Built a Neo4j knowledge graph linking job skills and candidate abilities.
- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation, and it became a reusable team best practice.
- Integrated Prometheus and Grafana for full-link observability, recognized by stakeholders and peers.
- Implemented RAG evidence retrieval with citation tracing for explainable conclusions.

Customer Service Agent (FastAPI + LangGraph)

- Designed tool orchestration and function calling to connect business systems, improving end-to-end delivery quality.
- Improved user satisfaction by 42%.
- Designed human fallback and escalation for better experience.
- Implemented intent recognition and slot filling to raise resolution rate, improving end-to-end delivery quality.

Highlights

- Established metrics and reviews around project management and planning to keep improving delivery quality
- Work around RESTful API design delivered about 42% efficiency gains and was reused across the team
- Continuously optimized root-cause analysis and postmortems, cutting key-path latency from 1800ms to 180ms
- Drove a focused effort on RAG retrieval and reranking, reducing related issues by about 38%
- Led a major technical effort involving Git collaboration and code review, staying stable at 12000 QPS peak

Skills

Core skills: Spring AI / FastAPI services, Agent, Spring AI, Prompt engineering and evaluation, Python
 Proficient in: Tool Orchestration, MCP, Observability, Milvus vector database, LLM observability and tracing

Familiar with: RAG, LangGraph, LLM, RAG retrieval and reranking, Milvus, LangGraph/LangChain agent orchestration

Self Evaluation

- Strong communication and collaboration with product, QA and ops to deliver projects on time
- Enjoys retrospectives, cares about code quality and observability, and makes data-driven decisions