

Ma Jiahao

Gender: Male Objective: AI Agent Engineer Location: Hangzhou

Phone: 13886843021 Email: majiahao6292@gmail.com GitHub: <https://github.com/majiahao403>

Education

2011.09 - 2015.06 Huazhong University of Science and Technology Electronic Information Engineering (M.S.)

Courses: Operating Systems, Computer Networks, Distributed Systems, Machine Learning, Design Patterns, Database Systems

Summary

Backend engineer building production AI-agent workflows, RAG pipelines and runtime observability, skilled in multi-agent orchestration and tool governance.

Work Experience

2018.07 - Present Alibaba Senior AI Agent Engineer

- Added structured constraints and validation on LLM output to reduce hallucination, raising key-field extraction accuracy above 90%, with experience documented and shared internally.
- Designed an agent runtime with route plan, tool budget, policy guardrails, trace replay and an evaluation feedback loop, and it became a reusable team best practice.
- Tuned Milvus index and sharding strategy, reducing retrieval latency from 1200ms to 380ms, significantly lowering maintenance cost.
- Handled 3380K agent invocations per day while keeping the online service stable, recognized by stakeholders and peers.
- Implemented hybrid RAG fusing BM25 lexical recall, vector search, reranking and fallback visibility, raising evidence hit rate by 48%, and related standards were adopted by the team.

2015.07 - 2018.06 Didi AI Agent Engineer

- Built backend of a multi-agent resume evaluation workflow chaining intent router, parser, JD matcher, technical/risk evaluators and report generator, supporting rapid business growth.
- Wired the full LLM call chain into Langfuse, tracking hit count, top score, fallback rate and latency per step.
- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 29%, and key metrics improved steadily.
- Built an offline evaluation set and automated scoring covering faithfulness, relevance and usability.

Projects

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Integrated Prometheus and Grafana for full-link observability.
- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation.
- Built a Neo4j knowledge graph linking job skills and candidate abilities, and it became a reusable team best practice.

- Implemented RAG evidence retrieval with citation tracing for explainable conclusions, and it became a reusable team best practice.

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Added reranking and citation tracing to reduce off-topic answers.
- Supported multi-turn conversation memory and context management.
- Implemented document chunking and multi-route recall to improve long-document QA, reused across multiple business lines.
- Built online evaluation and a feedback loop.

Customer Service Agent (FastAPI + LangGraph)

- Improved user satisfaction by 48%, significantly lowering maintenance cost.
- Designed tool orchestration and function calling to connect business systems, and related standards were adopted by the team.
- Implemented intent recognition and slot filling to raise resolution rate, reused across multiple business lines.
- Designed human fallback and escalation for better experience.

Highlights

- Led a major technical effort involving LLM observability and tracing, staying stable at 3000 QPS peak
- Work around Spring AI / FastAPI services delivered about 48% efficiency gains and was reused across the team
- Drove a focused effort on requirement analysis and documentation, reducing related issues by about 29%
- Continuously optimized unit testing and quality, cutting key-path latency from 1200ms to 380ms
- Built a reusable methodology and docs around Git collaboration and code review, adopted by several teams

Skills

Core skills: RAG retrieval and reranking, LangGraph, RAG, Spring AI, Milvus vector database

Proficient in: Prompt engineering and evaluation, LLM, Observability, Milvus, Spring AI / FastAPI services

Familiar with: MCP, Agent, Python, Tool Orchestration, LangGraph/LangChain agent orchestration, LLM observability and tracing

Self Evaluation

- Enjoys retrospectives, cares about code quality and observability, and makes data-driven decisions
- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end