

Ma Zihan

Gender: Female Objective: Data Platform Engineer Location: Wuhan

Phone: 19535952287 Email: mazihan1298@qq.com

## Education

2014.09 - 2018.06 Chongqing University Computer Science and Technology (B.S.)

Courses: Operating Systems, Computer Organization, Distributed Systems, Design Patterns, Compilers, Database Systems

GPA 3.3/4.0, top 15% in major; merit scholarship

## Summary

Data platform engineer focused on realtime warehouse, streaming ETL and batch processing with strong data modeling and governance.

## Work Experience

2020.07 - Present ByteDance Data Engineer

- Optimized Spark offline jobs, cutting runtime by 20% and resource cost by 51%.
- Tuned Kafka consumption and partitioning, significantly reducing backlog.
- Designed a layered warehouse (ODS/DWD/DWS/ADS) and data modeling for BI and ML features, with experience documented and shared internally.
- Governed metric definitions on a unified platform to reduce inconsistencies, recognized by stakeholders and peers.

2018.07 - 2020.06 Xiaohongshu Junior Data Engineer

- Processed 6110K records per day keeping T+1 and realtime paths stable.
- Built data-quality monitoring and lineage with automatic anomaly alerts, and it became a reusable team best practice.
- Supported realtime dashboards and risk features with second-level end-to-end latency, supporting rapid business growth.

## Projects

Realtime Data Platform (Flink + Kafka + ClickHouse)

- Processed 6110K events per day.
- Wrote to ClickHouse for realtime dashboards and ad-hoc queries.
- Processed Kafka realtime data with Flink at second-level end-to-end latency.
- Designed exactly-once semantics for data accuracy, recognized by stakeholders and peers.

Offline Warehouse & Metrics Platform (Spark + Hive)

- Built data-quality and lineage monitoring, significantly lowering maintenance cost.
- Built a layered warehouse with Spark and Hive.
- Unified metric definitions and dimensional modeling, significantly lowering maintenance cost.
- Optimized scheduling, cutting runtime by 20%.

## Highlights

- Led a major technical effort involving data analysis and metrics, staying stable at 3000 QPS peak
- Drove a focused effort on computer networking basics, reducing related issues by about 51%
- Built a reusable methodology and docs around Data warehouse modeling, adopted by several teams
- Work around Linux and shell basics delivered about 20% efficiency gains and was reused across the team
- Owned core modules that steadily served 6110K daily visits with 99.9%+ availability

#### Skills

Core skills: Flink stream processing, Data Governance, Spark, Data Warehouse

Proficient in: Python, Flink, Hive/ClickHouse, Spark batch processing

Familiar with: ETL, Data warehouse modeling, Kafka, Data governance and lineage, Hive, Kafka data ingestion

#### Self Evaluation

- Strong communication and collaboration with product, QA and ops to deliver projects on time
- Business-minded, good at balancing technical design against delivery timelines