

Wang Xue

Gender: Female Objective: Data Platform Engineer Location: Hangzhou

Phone: 15529495760 Email: wangxue9335@163.com GitHub: <https://github.com/wangxue428>

Education

2017.09 - 2021.06 Shandong University Computer Science and Technology (M.S.)

Courses: Machine Learning, Distributed Systems, Compilers, Database Systems, Data Structures and Algorithms, Operating Systems

GPA 3.4/4.0, top 5% in major; merit scholarship

Summary

Data platform engineer focused on realtime warehouse, streaming ETL and batch processing with strong data modeling and governance.

Work Experience

2023.07 - Present a SaaS company Data Engineer

- Governed metric definitions on a unified platform to reduce inconsistencies, reused across multiple business lines.
- Owned realtime warehouse and streaming ETL with Flink over massive Kafka event data.
- Supported realtime dashboards and risk features with second-level end-to-end latency.
- Designed a layered warehouse (ODS/DWD/DWS/ADS) and data modeling for BI and ML features, and key metrics improved steadily.

2021.07 - 2023.06 Xiaohongshu Junior Data Engineer

- Built data-quality monitoring and lineage with automatic anomaly alerts, supporting rapid business growth.
- Processed 7460K records per day keeping T+1 and realtime paths stable, and related standards were adopted by the team.
- Optimized Spark offline jobs, cutting runtime by 24% and resource cost by 16%, and related standards were adopted by the team.

Projects

Realtime Data Platform (Flink + Kafka + ClickHouse)

- Processed 7460K events per day, and key metrics improved steadily.
- Wrote to ClickHouse for realtime dashboards and ad-hoc queries.
- Processed Kafka realtime data with Flink at second-level end-to-end latency, and key metrics improved steadily.
- Designed exactly-once semantics for data accuracy.

Offline Warehouse & Metrics Platform (Spark + Hive)

- Optimized scheduling, cutting runtime by 24%, and it became a reusable team best practice.
- Unified metric definitions and dimensional modeling, recognized by stakeholders and peers.
- Built data-quality and lineage monitoring.
- Built a layered warehouse with Spark and Hive, significantly lowering maintenance cost.

Highlights

- Established metrics and reviews around design patterns and refactoring to keep improving delivery quality
- Continuously optimized computer networking basics, cutting key-path latency from 2200ms to 260ms
- Work around cross-team communication delivered about 24% efficiency gains and was reused across the team
- Led a major technical effort involving data analysis and metrics, staying stable at 30000 QPS peak
- Drove a focused effort on Flink stream processing, reducing related issues by about 16%

Skills

Core skills: Kafka data ingestion, Data governance and lineage, Spark batch processing, Spark

Proficient in: ETL, Data Warehouse, Kafka, Data Governance

Familiar with: Python, Hive/ClickHouse, Data warehouse modeling, Hive, Flink stream processing, Flink

Self Evaluation

- Stays calm under pressure and recovers services quickly during high-concurrency incidents
- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end