

Li Peng

Gender: Male Objective: Data Platform Engineer Location: Beijing

Phone: 18899247157 Email: lipeng6803@qq.com

Education

2015.09 - 2019.06 South China University of Technology Communication Engineering (B.S.)

Courses: Database Systems, Software Engineering, Machine Learning, Data Structures and Algorithms, Operating Systems, Distributed Systems

GPA 3.7/4.0, top 15% in major; merit scholarship

Summary

Data platform engineer focused on realtime warehouse, streaming ETL and batch processing with strong data modeling and governance.

Work Experience

2021.07 - Present Bilibili Data Engineer

- Owned realtime warehouse and streaming ETL with Flink over massive Kafka event data.
- Built data-quality monitoring and lineage with automatic anomaly alerts, and it became a reusable team best practice.
- Governed metric definitions on a unified platform to reduce inconsistencies, and related standards were adopted by the team.
- Supported realtime dashboards and risk features with second-level end-to-end latency, ensuring online service stability.

2019.07 - 2021.06 NIO Junior Data Engineer

- Optimized Spark offline jobs, cutting runtime by 26% and resource cost by 16%.
- Tuned Kafka consumption and partitioning, significantly reducing backlog, and it became a reusable team best practice.
- Designed a layered warehouse (ODS/DWD/DWS/ADS) and data modeling for BI and ML features, and key metrics improved steadily.

Projects

Offline Warehouse & Metrics Platform (Spark + Hive)

- Unified metric definitions and dimensional modeling.
- Optimized scheduling, cutting runtime by 26%.
- Built a layered warehouse with Spark and Hive.
- Built data-quality and lineage monitoring.

Realtime Data Platform (Flink + Kafka + ClickHouse)

- Processed Kafka realtime data with Flink at second-level end-to-end latency, recognized by stakeholders and peers.
- Designed exactly-once semantics for data accuracy, supporting rapid business growth.
- Wrote to ClickHouse for realtime dashboards and ad-hoc queries.
- Processed 1280K events per day.

Highlights

- Continuously optimized CI/CD pipelines, cutting key-path latency from 2200ms to 180ms
- Owned core modules that steadily served 1280K daily visits with 99.9%+ availability
- Led a major technical effort involving Spark batch processing, staying stable at 30000 QPS peak
- Drove a focused effort on cross-team communication, reducing related issues by about 16%
- Work around Kafka data ingestion delivered about 26% efficiency gains and was reused across the team

Skills

Core skills: Data Warehouse, Spark, Data warehouse modeling, Data Governance

Proficient in: Python, ETL, Hive, Hive/ClickHouse

Familiar with: Flink, Flink stream processing, Kafka data ingestion, Kafka, Spark batch processing, Data governance and lineage

Self Evaluation

- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end
- Curious about new tech, continuously learning LLM and cloud-native topics and applying them in practice