Hu Li
Gender: Female   Objective: Data Platform Engineer   Location: Xiamen
Phone: 15625932865   Email: huli1962@outlook.com   GitHub: https://github.com/huli114

Education

2016.09 - 2020.06   Hunan University   Information Security (B.S.)
Courses: Database Systems, Machine Learning, Operating Systems, Computer Organization, Compilers, Distributed Systems
GPA 3.7/4.0, top 5% in major; merit scholarship

Summary

Data platform engineer focused on realtime warehouse, streaming ETL and batch processing with strong data modeling and governance.

Work Experience

2022.07 - Present   JD.com   Data Engineer
- Processed 8110K records per day keeping T+1 and realtime paths stable.
- Designed a layered warehouse (ODS/DWD/DWS/ADS) and data modeling for BI and ML features.
- Governed metric definitions on a unified platform to reduce inconsistencies.
- Supported realtime dashboards and risk features with second-level end-to-end latency, and related standards were adopted by the team.

2020.07 - 2022.06   Kuaishou   Junior Data Engineer
- Built data-quality monitoring and lineage with automatic anomaly alerts.
- Optimized Spark offline jobs, cutting runtime by 58% and resource cost by 25%, significantly lowering maintenance cost.
- Owned realtime warehouse and streaming ETL with Flink over massive Kafka event data, ensuring online service stability.

Projects

Offline Warehouse & Metrics Platform (Spark + Hive)
- Unified metric definitions and dimensional modeling.
- Built a layered warehouse with Spark and Hive.
- Optimized scheduling, cutting runtime by 58%, reused across multiple business lines.
- Built data-quality and lineage monitoring, significantly lowering maintenance cost.

Realtime Data Platform (Flink + Kafka + ClickHouse)
- Designed exactly-once semantics for data accuracy.
- Processed 8110K events per day, and key metrics improved steadily.
- Processed Kafka realtime data with Flink at second-level end-to-end latency.
- Wrote to ClickHouse for realtime dashboards and ad-hoc queries, recognized by stakeholders and peers.

Highlights

- Established metrics and reviews around cross-team communication to keep improving delivery quality

- Built a reusable methodology and docs around Linux and shell basics, adopted by several teams
- Drove a focused effort on unit testing and quality, reducing related issues by about 25%
- Work around Flink stream processing delivered about 58% efficiency gains and was reused across the team
- Led a major technical effort involving data analysis and metrics, staying stable at 20000 QPS peak

Skills

Core skills: Hive/ClickHouse, Flink, Data Warehouse, Hive

Proficient in: ETL, Data warehouse modeling, Data governance and lineage, Flink stream processing

Familiar with: Data Governance, Spark, Python, Kafka data ingestion, Kafka, Spark batch processing

Self Evaluation

- Curious about new tech, continuously learning LLM and cloud-native topics and applying them in practice
- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end