

Dong Yuxin

Gender: Female Objective: LLM Application Engineer Location: Guangzhou

Phone: 18300099843 Email: dongyuxin7120@163.com GitHub: <https://github.com/dongyuxin122>

Education

2014.09 - 2018.06 Beijing Institute of Technology Network Engineering (B.S.)

Courses: Machine Learning, Computer Networks, Computer Organization, Operating Systems, Data Structures and Algorithms, Distributed Systems

GPA 3.5/4.0, top 10% in major; merit scholarship

Summary

LLM application engineer focused on RAG retrieval, vector store optimization, reranking and offline evaluation.

Work Experience

2020.07 - Present Li Auto LLM Application Engineer

- Added structured constraints and validation on LLM output to reduce hallucination, raising key-field extraction accuracy above 90%, and related standards were adopted by the team.
- Integrated an MCP resume-evidence service and dynamic skill loading for evidence synthesis and tool governance, supporting rapid business growth.
- Handled 5850K agent invocations per day while keeping the online service stable, ensuring online service stability.
- Built an offline evaluation set and automated scoring covering faithfulness, relevance and usability, and it became a reusable team best practice.

2018.07 - 2020.06 Kuaishou Junior LLM Application Engineer

- Tuned Milvus index and sharding strategy, reducing retrieval latency from 1800ms to 180ms, and key metrics improved steadily.
- Built backend of a multi-agent resume evaluation workflow chaining intent router, parser, JD matcher, technical/risk evaluators and report generator, improving end-to-end delivery quality.
- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 30%, ensuring online service stability.

Projects

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Added reranking and citation tracing to reduce off-topic answers, ensuring online service stability.
- Built online evaluation and a feedback loop, ensuring online service stability.
- Implemented document chunking and multi-route recall to improve long-document QA, and it became a reusable team best practice.
- Supported multi-turn conversation memory and context management.

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Integrated Prometheus and Grafana for full-link observability.
- Implemented RAG evidence retrieval with citation tracing for explainable conclusions.

- Built a Neo4j knowledge graph linking job skills and candidate abilities, ensuring online service stability.
- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation, with experience documented and shared internally.

Customer Service Agent (FastAPI + LangGraph)

- Implemented intent recognition and slot filling to raise resolution rate, improving end-to-end delivery quality.
- Designed human fallback and escalation for better experience, improving end-to-end delivery quality.
- Improved user satisfaction by 26%, supporting rapid business growth.
- Designed tool orchestration and function calling to connect business systems, and it became a reusable team best practice.

Highlights

- Established metrics and reviews around cross-team communication to keep improving delivery quality
- Owned core modules that steadily served 5850K daily visits with 99.9%+ availability
- Led a major technical effort involving Git collaboration and code review, staying stable at 3000 QPS peak
- Work around RESTful API design delivered about 26% efficiency gains and was reused across the team
- Drove a focused effort on project management and planning, reducing related issues by about 30%

Skills

Core skills: LangChain, Prompt Engineering, LangGraph/LangChain agent orchestration, Spring AI / FastAPI services, RAG retrieval and reranking

Proficient in: Rerank, LLM observability and tracing, Evaluation, Milvus, Embedding

Familiar with: Vector Search, RAG, Python, Milvus vector database, Knowledge Base, Prompt engineering and evaluation

Self Evaluation

- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end
- Curious about new tech, continuously learning LLM and cloud-native topics and applying them in practice