

Chen Zihan

Gender: Female Objective: LLM Application Engineer Location: Xian

Phone: 18860835701 Email: chenzihan326@outlook.com

Education

2016.09 - 2020.06 University of Electronic Science and Technology of China Data Science and Big Data Technology (B.S.)

Courses: Software Engineering, Data Structures and Algorithms, Database Systems, Machine Learning, Computer Organization, Computer Networks

GPA 3.6/4.0, top 5% in major; merit scholarship

Summary

LLM application engineer focused on RAG retrieval, vector store optimization, reranking and offline evaluation.

Work Experience

2022.07 - Present Tencent LLM Application Engineer

- Built backend of a multi-agent resume evaluation workflow chaining intent router, parser, JD matcher, technical/risk evaluators and report generator, significantly lowering maintenance cost.
- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 41%, supporting rapid business growth.
- Designed an agent runtime with route plan, tool budget, policy guardrails, trace replay and an evaluation on feedback loop, ensuring online service stability.
- Wired the full LLM call chain into Langfuse, tracking hit count, top score, fallback rate and latency per step, and related standards were adopted by the team.

2020.07 - 2022.06 NetEase Junior LLM Application Engineer

- Integrated an MCP resume-evidence service and dynamic skill loading for evidence synthesis and tool governance.
- Handled 3550K agent invocations per day while keeping the online service stable, improving end-to-end delivery quality.
- Tuned Milvus index and sharding strategy, reducing retrieval latency from 1200ms to 180ms, and key metrics improved steadily.

Projects

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Implemented RAG evidence retrieval with citation tracing for explainable conclusions.
- Built a Neo4j knowledge graph linking job skills and candidate abilities, recognized by stakeholders and peers.
- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation, and key metrics improved steadily.
- Integrated Prometheus and Grafana for full-link observability, recognized by stakeholders and peers.

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Implemented document chunking and multi-route recall to improve long-document QA.
- Built online evaluation and a feedback loop.
- Added reranking and citation tracing to reduce off-topic answers, and key metrics improved steadily.
- Supported multi-turn conversation memory and context management.

Customer Service Agent (FastAPI + LangGraph)

- Implemented intent recognition and slot filling to raise resolution rate, ensuring online service stability.
- Designed tool orchestration and function calling to connect business systems, significantly lowering maintenance cost.
- Improved user satisfaction by 35%, and it became a reusable team best practice.
- Designed human fallback and escalation for better experience, with experience documented and shared internally.

Highlights

- Drove a focused effort on unit testing and quality, reducing related issues by about 41%
- Established metrics and reviews around Spring AI / FastAPI services to keep improving delivery quality
- Owned core modules that steadily served 3550K daily visits with 99.9%+ availability
- Built a reusable methodology and docs around RAG retrieval and reranking, adopted by several teams
- Continuously optimized cross-team communication, cutting key-path latency from 1200ms to 180ms

Skills

Core skills: LangChain, Knowledge Base, Milvus vector database, Python, Vector Search

Proficient in: RAG retrieval and reranking, Evaluation, Rerank, Prompt Engineering, LangGraph/LangChain agent orchestration

Familiar with: RAG, LLM observability and tracing, Prompt engineering and evaluation, Spring AI / FastAPI services, Milvus, Embedding

Self Evaluation

- Solid engineering fundamentals and strong troubleshooting skills, able to own a module end to end
- Strong communication and collaboration with product, QA and ops to deliver projects on time