

Dong Yuhang

Gender: Male Objective: LLM Application Engineer Location: Hangzhou

Phone: 17581956572 Email: dongyuhang6056@foxmail.com

Education

2016.09 - 2020.06 University of Electronic Science and Technology of China Electronic Information Engineering (M.S.)

Courses: Operating Systems, Compilers, Machine Learning, Computer Networks, Computer Organization, Data Structures and Algorithms

Summary

LLM application engineer focused on RAG retrieval, vector store optimization, reranking and offline evaluation.

Work Experience

2022.07 - Present Tencent LLM Application Engineer

- Integrated an MCP resume-evidence service and dynamic skill loading for evidence synthesis and tool governance, with experience documented and shared internally.
- Designed an agent runtime with route plan, tool budget, policy guardrails, trace replay and an evaluation on feedback loop.
- Built an offline evaluation set and automated scoring covering faithfulness, relevance and usability, supporting rapid business growth.
- Built backend of a multi-agent resume evaluation workflow chaining intent router, parser, JD matcher, technical/risk evaluators and report generator, and related standards were adopted by the team.

2020.07 - 2022.06 a fintech company Junior LLM Application Engineer

- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 30%.
- Added structured constraints and validation on LLM output to reduce hallucination, raising key-field extraction accuracy above 90%, significantly lowering maintenance cost.
- Implemented hybrid RAG fusing BM25 lexical recall, vector search, reranking and fallback visibility, raising evidence hit rate by 43%.

Projects

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Implemented document chunking and multi-route recall to improve long-document QA, improving end-to-end delivery quality.
- Supported multi-turn conversation memory and context management, supporting rapid business growth.
- Built online evaluation and a feedback loop, supporting rapid business growth.
- Added reranking and citation tracing to reduce off-topic answers.

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Built a Neo4j knowledge graph linking job skills and candidate abilities.
- Integrated Prometheus and Grafana for full-link observability.

- Implemented RAG evidence retrieval with citation tracing for explainable conclusions.
- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation.

Customer Service Agent (FastAPI + LangGraph)

- Designed human fallback and escalation for better experience.
- Designed tool orchestration and function calling to connect business systems.
- Improved user satisfaction by 43%, supporting rapid business growth.
- Implemented intent recognition and slot filling to raise resolution rate, recognized by stakeholders and peers.

Highlights

- Owned core modules that steadily served 6570K daily visits with 99.9%+ availability
- Built a reusable methodology and docs around RAG retrieval and reranking, adopted by several teams
- Work around cross-team communication delivered about 43% efficiency gains and was reused across the team
- Drove a focused effort on Milvus vector database, reducing related issues by about 30%
- Established metrics and reviews around Linux and shell basics to keep improving delivery quality

Skills

Core skills: Python, RAG, Milvus vector database, Prompt engineering and evaluation, Knowledge Base

Proficient in: Spring AI / FastAPI services, RAG retrieval and reranking, Prompt Engineering, LangGraph/LangChain agent orchestration, Milvus

Familiar with: Vector Search, LLM observability and tracing, Embedding, LangChain, Rerank, Evaluation

Self Evaluation

- Curious about new tech, continuously learning LLM and cloud-native topics and applying them in practice
- Strong communication and collaboration with product, QA and ops to deliver projects on time