

Liang Jun

Gender: Male Objective: LLM Application Engineer Location: Chengdu

Phone: 19909841286 Email: liangjun5915@qq.com GitHub: <https://github.com/liangjun828>

Education

2014.09 - 2018.06 Southeast University Electronic Information Engineering (M.S.)

Courses: Database Systems, Computer Organization, Operating Systems, Machine Learning, Software Engineering, Data Structures and Algorithms

Summary

LLM application engineer focused on RAG retrieval, vector store optimization, reranking and offline evaluation.

Work Experience

2020.07 - Present Alibaba LLM Application Engineer

- Implemented hybrid RAG fusing BM25 lexical recall, vector search, reranking and fallback visibility, raising evidence hit rate by 53%, and key metrics improved steadily.
- Integrated an MCP resume-evidence service and dynamic skill loading for evidence synthesis and tool governance, with experience documented and shared internally.
- Built an offline evaluation set and automated scoring covering faithfulness, relevance and usability, reused across multiple business lines.
- Added structured constraints and validation on LLM output to reduce hallucination, raising key-field extraction accuracy above 90%.

2018.07 - 2020.06 NetEase Junior LLM Application Engineer

- Tuned Milvus index and sharding strategy, reducing retrieval latency from 1800ms to 320ms, and it became a reusable team best practice.
- Handled 8260K agent invocations per day while keeping the online service stable, and related standards were adopted by the team.
- Designed prompt templates with version management supporting gray release and A/B, improving iteration speed by 22%, and related standards were adopted by the team.

Projects

Customer Service Agent (FastAPI + LangGraph)

- Designed tool orchestration and function calling to connect business systems.
- Designed human fallback and escalation for better experience.
- Implemented intent recognition and slot filling to raise resolution rate.
- Improved user satisfaction by 53%, with experience documented and shared internally.

ResumAI Agent - Intelligent Resume Evaluation Platform (Spring Boot + Vue3 + Milvus + Neo4j + DeepSeek)

- Built a Neo4j knowledge graph linking job skills and candidate abilities, and key metrics improved steadily.
- Integrated Prometheus and Grafana for full-link observability, significantly lowering maintenance cost.
- Implemented RAG evidence retrieval with citation tracing for explainable conclusions, and it became a r

reusable team best practice.

- Orchestrated a multi-agent evaluation flow with a DAG covering parsing, matching, evaluation and report generation, reused across multiple business lines.

Enterprise Knowledge Base QA (LangChain + Milvus + Redis)

- Implemented document chunking and multi-route recall to improve long-document QA, and key metrics improved steadily.
- Supported multi-turn conversation memory and context management.
- Built online evaluation and a feedback loop, with experience documented and shared internally.
- Added reranking and citation tracing to reduce off-topic answers.

Highlights

- Continuously optimized data analysis and metrics, cutting key-path latency from 1800ms to 320ms
- Led a major technical effort involving cross-team communication, staying stable at 8000 QPS peak
- Work around LLM observability and tracing delivered about 53% efficiency gains and was reused across the team
- Drove a focused effort on CI/CD pipelines, reducing related issues by about 22%
- Established metrics and reviews around project management and planning to keep improving delivery quality

Skills

Core skills: Prompt engineering and evaluation, LLM observability and tracing, RAG, Milvus, Spring AI / FastAPI services

Proficient in: LangChain, LangGraph/LangChain agent orchestration, Evaluation, Python, Rerank

Familiar with: Knowledge Base, Prompt Engineering, Milvus vector database, RAG retrieval and reranking, Vector Search, Embedding

Self Evaluation

- Curious about new tech, continuously learning LLM and cloud-native topics and applying them in practice
- Enjoys retrospectives, cares about code quality and observability, and makes data-driven decisions